

Class Documentation for Inference Algorithms

Contents

1	Overview	2
2	Prior Parameter Classes	3
2.1	GMPriorParameters Class	3
2.1.1	Description	3
2.1.2	Constructor	3
2.2	BMPriorParameters Class	3
2.2.1	Description	3
2.2.2	Constructor	4
3	Gibbs Sampler	5
3.1	GibbsModel Class	5
3.1.1	Description	5
3.1.2	Methods	5
3.2	BMMGibbs Class	6
3.2.1	Description	6
3.2.2	Constructor	6
3.2.3	Methods	6
3.3	GMMGibbs Class	7
3.3.1	Description	7
3.3.2	Constructor	7
3.3.3	Methods	7
4	Variational Bayesian EM Algorithm (VBEM)	9
4.1	BMMVBEM Class	9
4.1.1	Description	9
4.1.2	Constructor	9
4.1.3	Methods	9
4.2	GMMVBEM Class	10
4.2.1	Description	10
4.2.2	Constructor	10
4.2.3	Methods	10
5	Expectation Maximization Classes	12
5.1	BMMEM Class	12
5.1.1	Description	12
5.1.2	Constructor	12
5.1.3	Methods	12
5.2	GMMEM Class	14
5.2.1	Description	14
5.2.2	Constructor	14
5.2.3	Methods	14

Chapter 1

Overview

All models have been implemented using an object-oriented approach, with a shared set of core methods across all models. The common methods include:

- **fit()** : This method executes the training process and updates the model's learned parameters. It takes as input a NumPy array of shape $(N \times D)$, where N is the number of data points and D is the number of features. Additional optional arguments specific to each model may also be included and are detailed below.
- **predict()** : This method accepts a NumPy array of shape $(N \times D)$ as input and returns an array of component assignments corresponding to the input data points. Each model may include additional optional arguments, which are described in the following sections.
- **impute()** : This method takes an input data matrix with missing values and returns a new matrix where the missing values have been imputed. The imputation process is model-specific and can handle different approaches based on the model used.

In the documentation for each specific model below, we only include user-facing methods and have omitted internal methods.

The models in this documentation are as follows:

1. **GMMGibbs** : Gibbs sampler for GMM
2. **BMMGibbs** : Gibbs sampler for BMM
3. **GMMVBEM** : VBEM for GMM
4. **BMMVBEM** : VBEM for BMM
5. **GMMEM** : EM for GMM
6. **BMMEM** : EM for BMM

This also includes hyper-parameter classes used to instantiate prior distribution hyper-parameter values for each inference class

Chapter 2

Prior Parameter Classes

2.1 GMPriorParameters Class

2.1.1 Description

`GMPriorParameters` defines the prior parameters for a Gaussian Mixture Model (GMM). These parameters include the prior distributions for the mixture components, component means, covariances, and other hyperparameters required for fitting the model.

2.1.2 Constructor

```
def init(self, data, K):
```

Parameters:

- `data`: The data matrix (ND), where N is the number of data points and D is the number of features.
- `K`: The number of components in the Gaussian Mixture Model (GMM).

Attributes:

- `K`: The number of components in the mixture model.
- `α_0`: The Dirichlet prior for the component weights (K).
- `m_0`: The prior means for the components ($K \times D$).
- `S_0`: The prior covariance matrices for the components ($K \times D \times D$).
- `k_0`: The prior scaling factor for the component means (K).
- `ν_0`: The prior degrees of freedom for the component covariances (K).
- `o, h`: Hyperparameters for MNAR data handling using the Beta distribution ($K \times D$).

2.2 BMMPriorParameters Class

2.2.1 Description

`BMMPriorParameters` defines the prior parameters for a Bernoulli Mixture Model (BMM). Similar to the `GMPriorParameters` class, this class defines priors for the mixture components, component-specific parameters, and missing data handling.

2.2.2 Constructor

```
def init(self, data, K):
```

Parameters:

- **data**: The data matrix (ND), where N is the number of data points and D is the number of features.
- **K**: The number of components in the Bernoulli Mixture Model (BMM).

Attributes:

- **K**: The number of components in the mixture model.
- **α_0** : The Dirichlet prior for the component weights (K).
- **a_0**: The prior alpha parameters for the Bernoulli distribution ($K \times D$).
- **b_0**: The prior beta parameters for the Bernoulli distribution ($K \times D$).
- **o, h**: Hyperparameters for missing data handling using the Beta distribution ($K \times D$).

Chapter 3

Gibbs Sampler

3.1 GibbsModel Class

3.1.1 Description

`GibbsModel` is an abstract base class for `BMMGibbs` and `GMMGibbs` models. Thus the methods listed below are shared among the two BMM and GMM specific Gibbs sampling classes.

3.1.2 Methods

`fit`

```
@abstractmethod
def fit(self, X, num_iters=6000, burn=2000, mmar=False, collapse=False):
```

This method must be implemented by subclasses to perform the Gibbs sampling algorithm. It fits the model to the data and stores the posterior samples.

Parameters:

- `X`: The input data matrix, possibly containing missing values.
- `num_iters`: The number of iterations to run the Gibbs sampler.
- `burn`: The number of burn-in iterations to discard.
- `mmar`: Boolean indicating whether to model MNAR (Missing Not At Random) missingness.
- `collapse`: Boolean indicating whether to collapse the model.

`get_map_params`

```
def get_map_params(self):
```

Returns the parameters corresponding to the maximum a posteriori (MAP) estimate.

`get_summarizing_results`

```
def get_summarizing_results(self, y):
```

Computes the average and standard deviation for metrics such as Adjusted Rand Index (ARI), log-likelihood, and posterior likelihood. **Parameters:**

- `y`: ground truth labels corresponding to sample used to fit model

`get_aligned_param_means`

```
def get_aligned_param_means(self):
```

Returns the mean of the aligned parameters across all samples.

3.2 BMMGibbs Class

3.2.1 Description

BMMGibbs implements a Bernoulli Mixture Model (BMM) with missing data using Gibbs Sampling. The class is derived from `GibbsModel` and includes specific methods for handling binary data and missing values.

3.2.2 Constructor

```
def __init__(self, priorParameters: BMMPriorParameters):
```

Parameters:

- `priorParameters`: An instance of the prior distribution for the Bernoulli mixture model.

3.2.3 Methods

`log_likelihood`

```
def log_likelihood(self, X_new, mnar=False):
```

Computes the log-likelihood of given input data matrix using the current model parameters. **Parameters:**

- `X_new`: data to compute likelihood using current parameters
- `mnar` : boolean indicating whether to use MNAR approach

Returns:

- The total log-likelihood as a scalar

`fit`

```
def fit(self, X, num_iters=6000, burn=2000, mnar=False, collapse=False):
```

Performs the Gibbs sampling for the Bernoulli mixture model, including missing data imputation and parameter updates.

Parameters:

- `X`: data to fit model on
- `num_iters` : number of iterations to run for Gibbs
- `burn` : number of samples to burn
- `mnar` : boolean indicating whether to use MNAR approach
- `collapse` : boolean indicating whether to use collapsed approach

impute

```
def impute(self, X_new, sample=None, eps=1e-14):
```

Imputes missing values in the data using the fitted model parameters.

Parameters:

- `X_new`: the sample to impute
- `sample` : allows specifying what Gibbs sample to use

Returns: the input data matrix with missing entries imputed

predict

```
def predict(self, X_new):
```

Uses fitted parameters to predict component assignments for a given input

Parameters:

- `X_new`: The input data matrix with or without missing values.
- `mnr` : boolean indicating whether to use MNAR approach

Returns:

- Numpy array of component assignments ($N,$)

3.3 GMMGibbs Class

3.3.1 Description

GMMGibbs implements a Gaussian Mixture Model (GMM) with missing data using Gibbs Sampling. This class is derived from `GibbsModel` and uses multivariate normal distributions for modeling continuous data.

3.3.2 Constructor

```
def __init__(self, priorParameters: GMPriorParameters):
```

Parameters:

- `priorParameters`: An instance of the prior distribution for the Gaussian mixture model.

3.3.3 Methods

log_likelihood

```
def log_likelihood(self, X_new, mnr=False):
```

Computes the log-likelihood of given input data matrix using the current model parameters. **Parameters:**

- `X_new`: data to compute likelihood using current parameters

- `mnar` : boolean indicating whether to use MNAR approach

Returns:The total log-likelihood as a scalar

fit

```
def fit(self, X, num_iters=6000, burn=2000, mnar=False, collapse=False):
```

Performs Gibbs sampling for the Gaussian mixture model, including parameter updates and missing data imputation.

Parameters:

- `mnar`: boolean indicating whether to use MNAR generative approach
- `collapse` : boolean indicator whether to use collapsed Gibbs

impute

```
def impute(self, X_new, sample=None, eps=1e-14):
```

Imputes missing values in the data using the Gaussian mixture model parameters.

Parameters:

- `X_new`: the sample to impute
- `sample` : allows specifying what Gibbs sample to use

Returns: the input data matrix with missing entries imputed

predict

```
def predict(self, X_new):
```

Uses fitted parameters to predict component assignments for a given input

Parameters:

- `X_new`: The input data matrix with or without missing values.
- `mnar` : boolean indicating whether to use MNAR approach

Returns:

- Numpy array of component assignments ($N,$)

Chapter 4

Variational Bayesian EM Algorithm (VBEM)

4.1 BMMVBEM Class

4.1.1 Description

BMMVBEM implements a Bernoulli Mixture Model (BMM) with missing data using VBEM. It is derived from the `VBEMModel` abstract class and includes specific methods for handling binary data with missing values using variational inference.

4.1.2 Constructor

```
def init(self, priorParameters : BMMPriorParameters):
```

Parameters:

- **priorParameters:** An instance of the prior distribution for the Bernoulli mixture model.

4.1.3 Methods

fit

```
def fit(self, X, mode=0, max_iters=200, tol=1e-4):
```

Performs the VBEM algorithm for the Bernoulli mixture model. The algorithm iteratively updates the parameters until convergence. **Parameters:**

- **X:** The input data matrix (N x D).
- **mode:** Determines the update mode for π . Mode 0 uses variational updates, and mode 1 uses MLE-style updates.
- **max_iters:** The maximum number of iterations to run.
- **tol:** The convergence tolerance for the ELBO.

impute

```
def impute(self, X_new, eps=1e-14):
```

Imputes missing values in the data using the fitted model parameters.

Parameters:

- `X_new`: The input data matrix with missing values.
- `eps`: A small constant to prevent division by zero.

Returns: the input data matrix with missing entries imputed

predict

```
def predict(self, X_new):
```

Uses fitted parameters to predict component assignments for a given input

Parameters:

- `X_new`: The input data matrix with or without missing values.

Returns:

- Numpy array of component assignments ($N,$)

4.2 GMMVBEM Class

4.2.1 Description

GMMVBEM implements a Gaussian Mixture Model (GMM) with missing data using VBEM. It is derived from the `VBEMModel` abstract class and uses multivariate normal distributions to model continuous data.

4.2.2 Constructor

```
def init(self, priorParameters : GMPriorParameters):
```

Parameters:

- `priorParameters`: An instance of the prior distribution for the Gaussian mixture model.

4.2.3 Methods

fit

```
def fit(self, X, max_iters=200, tol=1e-4):
```

Performs the VBEM algorithm for the Gaussian mixture model

Parameters:

- `X`: The input data matrix ($N \times D$).
- `max_iters`: The maximum number of iterations to run.
- `tol`: The convergence tolerance for the ELBO.

impute

```
def impute(self, X_new, eps=1e-14):
```

Imputes missing values in the data using the learned variational posteriors.

Parameters:

- **X_new**: The input data matrix with missing values.
- **eps**: A small constant to prevent division by zero.

Returns: the input data matrix with missing entries imputed

predict

```
def predict(self, X_new):
```

Uses fitted parameters to predict component assignments for a given input

Parameters:

- **X_new**: The input data matrix with or without missing values.

Returns:

- Numpy array of component assignments ($N,$)

Chapter 5

Expectation Maximization Classes

5.1 BMMEM Class

5.1.1 Description

BMMEM implements the Expectation-Maximization (EM) algorithm for Bernoulli Mixture Models (BMM).

5.1.2 Constructor

```
def init(self, K, complete_case=False):
```

Parameters:

- **K**: The number of components in the mixture model.

Attributes:

- **rng**: Random number generator for reproducibility.
- **K**: The number of components in the mixture model.
- **fitted**: A boolean indicating if the model has been fitted.
- **complete_case**: Whether to use complete case analysis for missing data.

5.1.3 Methods

fit

```
def fit(self, X, max_iters=200, tol=1e-4):
```

The `fit` method runs the EM algorithm to fit the BMM to the data. It iterates between the E-step and M-step until convergence, based on the provided tolerance. It also handles missing data using the `complete_case` flag.

Parameters:

- **X**: The input data matrix (N x D), where N is the number of data points and D is the number of features.
- **max_iters**: The maximum number of iterations to run the EM algorithm.
- **tol**: The convergence tolerance for the log-likelihood change between iterations.

Returns:

- A dictionary containing the following keys:
 - \mathbf{z} : Cluster assignments for each data point (N).
 - $\boldsymbol{\pi}$: Mixture component weights (K).
 - $\boldsymbol{\theta}$: Component parameters (K x D).
 - `loglike`: Log-likelihood values over iterations.

compute_responsibility

```
def compute_responsibility(self, X_new, eps=1e-14):
```

Computes the responsibility matrix R for a new set of data points, which indicates the probability that each data point belongs to each component of the mixture model. This method handles missing data according to the `complete_case` argument.

Parameters:

- `X_new`: The input data matrix (N x D) with missing values (represented as NaN).
- `eps`: A small constant to prevent numerical issues.

Returns:

- R : The responsibility matrix (N x K), with each row summing to 1.
- `loglike`: The log-likelihood of the data given the current model parameters.

impute

```
def impute(self, X_new, eps=1e-14):
```

Imputes missing values in the data using the fitted BMM parameters. The imputation is done based on the responsibility matrix R and the model parameters $\boldsymbol{\theta}$ for each mixture component.

Parameters:

- `X_new`: The input data matrix (N x D) with missing values (NaN).
- `eps`: A small constant to prevent numerical issues during the imputation process.

Returns:

- `X_filled`: The imputed data matrix (N x D), with missing values replaced by imputed values.

log_likelihood

```
def log_likelihood(self, X_new):
```

Computes the log-likelihood of a new set of data points under the current model parameters. This is used for model evaluation and convergence checking.

Parameters:

- `X_new`: The input data matrix (N x D).

Returns:

- `loglike`: The log-likelihood of the data under the current model.

predict

```
def predict(self, X_new):
```

Predicts the cluster assignments for a new set of data points.

Parameters:

- **X_new**: The input data matrix (N x D) with missing values.

Returns:

- A vector of cluster assignments (N), indicating the predicted mixture component for each data point.

5.2 GMMEM Class

5.2.1 Description

GMMEM implements the Expectation-Maximization (EM) algorithm for Gaussian Mixture Models (GMM) with missing data.

5.2.2 Constructor

```
def init(self, K, complete_case=False):
```

Parameters:

- **K**: The number of components in the mixture model.
- **complete_case**: Boolean indicating whether to perform complete case analysis (i.e., ignore incomplete data points).

Attributes:

- **rng**: Random number generator for reproducibility.
- **K**: The number of components in the mixture model.
- **fitted**: A boolean indicating if the model has been fitted.
- **complete_case**: Whether to use complete case analysis for missing data.

5.2.3 Methods

fit

```
def fit(self, X, max_iters=200, tol=1e-4):
```

The `fit` method runs the EM algorithm to fit the GMM to the data. It iterates between the E-step and M-step until convergence, based on the provided tolerance. It also handles missing data using the `complete_case` flag.

Parameters:

- **X**: The input data matrix (N x D), where N is the number of data points and D is the number of features.
- **max_iters**: The maximum number of iterations to run the EM algorithm.

- `tol`: The convergence tolerance for the log-likelihood change between iterations.

Returns:

- A dictionary containing the following keys:
 - `z`: Cluster assignments for each data point (N).
 - `π` : Mixture component weights (K).
 - `μ` : Component means (K x D).
 - `Σ` : Component covariances (K x D x D).
 - `loglike`: Log-likelihood values over iterations.

compute_responsibility

```
def compute_responsibility(self, X_new, eps=1e-10):
```

Computes the responsibility matrix R for a new set of data points. This method handles missing data according to the `complete_case` argument.

Parameters:

- `X_new`: The input data matrix (N x D) with missing values (represented as NaN).
- `eps`: A small constant to prevent numerical issues.

Returns:

- `R`: The responsibility matrix (N x K), with each row summing to 1.
- `loglike`: The log-likelihood of the data given the current model parameters.

impute

```
def impute(self, X_new, eps=1e-14):
```

Imputes missing values in the data using the fitted GMM parameters. The imputation is done based on the responsibility matrix R and the model parameters.

Parameters:

- `X_new`: The input data matrix (N x D) with missing values (NaN).
- `eps`: A small constant to prevent numerical issues during the imputation process.

Returns:

- `X_filled`: The imputed data matrix (N x D), with missing values replaced by imputed values.

log_likelihood

```
def log_likelihood(self, X_new):
```

Computes the log-likelihood of a new set of data points under the current model parameters. This is used for model evaluation and convergence checking.

Parameters:

- **X_{new}**: The input data matrix ($N \times D$).

Returns:

- **loglike**: The log-likelihood of the data under the current model.

predict

```
def predict(self, Xnew):
```

Predicts the cluster assignments for a new set of data points based on the current responsibility matrix.

Parameters:

- **X_{new}**: The input data matrix ($N \times D$) with missing values.

Returns:

- A vector of cluster assignments (N), indicating the predicted mixture component for each data point.